

Variable clustering for Husler-Reiss graphical models

CAPEL Alexandre

2025-02-26

Contents

1	Introduction	1
2	Husler-Reiss graphical model	2
2.1	Definition of a Husler-Reiss model	3
2.2	Characterisation of a HRGM	3
3	Clusterpath for GGM	4
3.1	Gaussian Graphical Model	5
3.2	Clusterpath algorithm	6
4	Clusterpath adapted for HRGM	7
4.1	Maximum likelihood for graphical model	7
4.2	Adaptation of the expression	8
4.3	Computation of the derivative	9
4.4	Choice of a step for descent	11
4.5	Simulation study	11
5	Application on flight delay data	11

We want to use the graphical model tools for extreme value theory to build a new way of clustering variable, as done for the graphical models for Gaussian vector (Touw et al. 2024).

Let $V = \{1, \dots, d\}$.

1 Introduction

For a multivariate random variable, it can be useful to know the dependence structure between the components. Particularly, we can summarise the conditional dependence structure with a graph $\mathcal{G} = (V, E)$ with $E \subset V \times V$ as below :

For classical conditional independence (in term of density factorisation), we call such variables graphical models.




Figure 1: Exemple of a graph

Construction of a graph

Let X a graphical model according to the graph $\mathcal{G} = (V, E)$.

Then, by definition, there is the relation :

$$(i, j) \in E \iff X_i \perp\!\!\!\perp X_j \mid X_{V \setminus \{i, j\}}$$

It is the pairwise Markov property.

The graphoid structure of such a relation gives us the equivalence between pairwise Markov property and the global Markov property :

$$A \perp_{\mathcal{G}} B \mid C \implies X_A \perp\!\!\!\perp X_B \mid X_C$$

where $\perp_{\mathcal{G}}$ is the separation relation between sets of nodes.

We would like to cluster the variable using the graphical model structure to be able to get an interpretation of the clusters and then reduce the dimension of the graph. In that sense, we would have the nodes as the clusters and the edge thanks to the global Markov properties relationship existing between these.

However for general graph, it is not easy. Indeed even with three clusters, we can have this type of situation :

and each time, we have the fact that :

$$X_A \perp\!\!\!\perp X_B \mid X_C$$

So we want to :

- get a **unique** decomposition using the graphical model structure.
- link this decomposition to a **way of clustering the variables**.

2 Husler-Reiss graphical model

In this section, we will present quickly the Husler-Reiss distribution in the MGPD case, and a present the notion of conditional independence in this context, together with the characterisation of this conditional independence for Husler-Reiss graphical models.




Figure 2: 3 different way of clustering where A is in blue, B in green and C in red.

2.1 Definition of a Husler-Reiss model

Now we consider a MGPD random vector Y indexed by V .

The Husler-Reiss model is a MGPD parameterized by a symmetric conditionally negative definite matrix Γ with $\text{diag}(\Gamma) = 0$.

One knows that every MGPD is defined by an exponent measure Λ giving a correspondence between MGEV and MGPD (see (Rootzén and Tajvidi 2006)). For the case of the Husler-Reiss distribution, the exponent measure is absolutely continuous with respect to the Lebesgue measure on the cone $\mathcal{E} = \mathbb{R}_+^d \setminus \{0\}$ and its derivative is given by (Engelke et al. 2012) for any $k \in V$:

$$\lambda(y) = y_k^{-2} \prod_{i \neq k} y_i^{-1} \phi(\tilde{y}_{\setminus k}; \Sigma^{(k)}), \quad y \in \mathcal{E}$$

where $\phi(\cdot, \Sigma)$ is the density function of a gaussian vector centered with covariance matrix Σ , $\tilde{y}_{\setminus k} = (\log(y_i/y_k) + \Gamma_{ik}/2)_{i \in V}$ and :

$$\Sigma_{ij}^{(k)} = \frac{1}{2}(\Gamma_{ik} + \Gamma_{kj} - \Gamma_{ij}) \quad \text{for } i, j \neq k$$

which is obviously definite positive.

2.2 Characterisation of a HRGM

For extreme value theory, the notion of conditional independence is very complicated to defined. First, for max-stable distribution with continuous positive density, the notion of conditional independence is equivalent to the global independence of the variables (Papastathopoulos reference). Moreover, for the MGPD case, the random vector is not even defined in a product space which make impossible the use of conditional independence.

Hopefully, (Engelke and Hitz 2020) build a new notion of conditional independence, adapted to MGPD distribution and then permit us to make graphical model with this type of distribution.

Let A, B and C a partition of V . Then for MGPLD random vector Y indexed by V , we say that Y_A is conditionally independent of Y_B given Y_C if :

$$Y_A^k \perp\!\!\!\perp Y_B^k \mid Y_C^k, \quad \forall k \in V.$$

where Y^k is defined as the vector Y conditionally to the event $\{Y_k > 1\}$.

We note then $Y_A \perp\!\!\!\perp_e Y_B \mid Y_C$.

Moreover, in the same article, they give a first characterisation of the extremal conditional independence for Husler-Reiss models :

Proposition. For a Husler-Reiss graphical model (HRGM) Y with variogram Γ , then for all $i, j \in V$ and $k \in V$ we have :

$$Y_i \perp\!\!\!\perp_e Y_j \mid Y_{V \setminus \{i, j\}} \Leftrightarrow \begin{cases} \Theta_{ij}^{(k)} = 0, & \text{if } i, j \neq k, \\ \sum \Theta_{lj}^{(k)} = 0, & \text{if } i = k, j \neq k, \\ \sum \Theta_{il}^{(k)} = 0, & \text{if } i \neq k, j = k \end{cases}$$

where $\Theta^{(k)}$ is the precision matrix of $\Sigma^{(k)}$ (i.e $\Theta^{(k)} = (\Sigma^{(k)})^{-1}$).

In (Hentschel, Engelke, and Segers 2023), they build an extended precision matrix Θ which summarize all the information we need for the conditional independence relationship for the extremal graphical models, in that sense :

$$Y_i \perp\!\!\!\perp_e Y_j \mid Y_{V \setminus \{i, j\}} \iff \Theta_{ij} = 0.$$

This matrix can be obtain by using some applications (which are bijections) that garanties a form of unicity of the spectral representation.

Therefore, let's consider the following applications :

$$\sigma : \Gamma \mapsto \Pi_d(-\frac{1}{2}\Gamma)\Pi_d, \quad \theta : \Gamma \mapsto \sigma(\Gamma)^+$$

where the matrix A^+ is the general inverse of A , and Π_d the orthogonal projection matrix in the space $< \mathbb{1} >^\perp$.

They show in (Hentschel, Engelke, and Segers 2023) that the above applications are homeomorphisms between the set of the strictly conditionally negative definite variogram matrix \mathcal{D}_d ad the set of symmetric positive semi-definite matrix with kernel equal to $< \mathbb{1} >$, denoted by \mathcal{P}_d^1 . More they show that :

$$\sigma^{-1}(\Sigma) = \gamma(\Sigma), \quad \theta^{-1}(\Theta) = \gamma(\Theta^+),$$

where $\gamma(\Sigma) = \mathbb{1}\text{diag}(\Sigma)^T + \text{diag}(\Sigma)\mathbb{1}^T - 2\Sigma$.

3 Clusterpath for GGM

In this section, we will present the matrix structure we will use for the Husler-Reiss graphical model. More, we will present an algorithm to estimate this graphical structure.

3.1 Gaussian Graphical Model

In (Touw et al. 2024), they build a graphical model that we can use for clustering, in the case of Gaussian graphical model (GGM).

For the GGM, there exists an easy characterisation of the conditional independence which is similar to HRGM for the extreme. For a Gaussian graphical model X with covariance matrix $\tilde{\Sigma}$, we have :

$$X_i \perp\!\!\!\perp_e X_j \mid X_{V \setminus \{i,j\}} \iff \tilde{\Theta}_{ij} = 0,$$

where $\tilde{\Theta} = \tilde{\Sigma}^{-1}$, the precision matrix.

Let assume that the variable X can be grouped in K clusters $\{C_1, \dots, C_K\}$ with $p_k = |C_k|$.

The goal was to encouraging clustering of the graph by forcing the precision matrix to have a K blocks structure as follows :

$$\tilde{\Theta} = \begin{pmatrix} (a_{11} - r_{11})I_{p_1} & 0 & \dots & 0 \\ 0 & (a_{22} - r_{22})I_{p_2} & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & (a_{KK} - r_{KK})I_{p_K} \end{pmatrix} + \begin{pmatrix} r_{11}\mathbb{1}\mathbb{1}^t & r_{12}\mathbb{1}\mathbb{1}^t & \dots & r_{1K}\mathbb{1}\mathbb{1}^t \\ r_{21}\mathbb{1}\mathbb{1}^t & r_{22}\mathbb{1}\mathbb{1}^t & \dots & r_{2K}\mathbb{1}\mathbb{1}^t \\ \vdots & \vdots & \ddots & \vdots \\ r_{K1}\mathbb{1}\mathbb{1}^t & r_{K2}\mathbb{1}\mathbb{1}^t & \dots & r_{KK}\mathbb{1}\mathbb{1}^t \end{pmatrix},$$

where A is a $K \times K$ diagonal matrix, R a $K \times K$ symmetric matrix, I_p the $p \times p$ identity matrix.

We can then get this type of “graph factorisation” which is unique due to precision matrix structure :




Figure 3: Graph factorisation thanks to the precision matrix structure

Thus, with this factorisation, we build three clusters with an interpretation of conditional independence between them.

For the estimation of the precision matrix, they use the following likelihood :

$$L(\Theta) = -\log(|\Theta|) + \text{tr}(\Sigma\Theta),$$

where $\log(|\cdot|)$ is the logarithm of the determinant and $\text{tr}(\cdot)$ the trace.

To get the maximum likelihood, they use a convex penalty to get the unknown block structure of the precision matrix. Thus, we got this optimisation program :

$$\hat{\Theta} = \arg \min_{\Theta} \left[-\log(|\Theta|) + \text{tr}(\bar{\Sigma}\Theta) + \lambda \mathcal{P}(\Theta) \right], \quad s.t. \Theta^t = \Theta, \quad \Theta > 0$$

where $\bar{\Sigma}$ is an estimation of the covariance matrix $\tilde{\Sigma}$.

From now, as we have for $i, j \in C_k$ that $\theta_{.i} = \theta_{.j}$, we will note θ_{C_k} the vector of the precision matrix of the cluster C_k .

3.2 Clusterpath algorithm

In order to find the groups in precision matrix $\tilde{\Theta}$, we will use the clusterpath algorithm from (Hocking et al., n.d.).

For these convex optimisation programs, we impose to the penalty function to be of the form :

$$\mathcal{P}(\Theta) = \sum_{i < j} w_{ij} D(\theta_{.i}, \theta_{.j}),$$

where w_{ij} are some positive weights, and D a distance in \mathbb{R}^d .

The distance D

We can use a lot of distance :

- with the l^p norm for $p \in [1, \infty]$.
- in particular l^1 , l^2 and l^∞ .
- in (Touw et al. 2024) they use another distance defined as :

$$D(\theta_{.i}, \theta_{.j}) = \sqrt{(\theta_{ii} - \theta_{jj})^2 + \sum_{k \neq i, j} (\theta_{ik} - \theta_{jk})^2}$$

which “can be interpreted as a group lasso penalty”.

Choice of w_{ij}

The choice of w_{ij} is also free, even if they present one which seems better (or nearer from the data) using :

$$w_{ij} = \exp(-a \|\theta_{.i} - \theta_{.j}\|^2)$$

where $\|\cdot\|$ is the l^2 norm.

Clusterpath algorithm

The algorithm is a gradient descent algorithm, adding conditions to detect clusters and fuse variables.

Algorithm 1 : Clusterpath

Input: initial guess Θ , initial estimation $\bar{\Sigma}$, initial clusters, weight w_{ij} , regularisation λ
 $G \leftarrow \text{gradient}(\cdot)$
while $\|G\| > \varepsilon$ **do**
 $\Theta \leftarrow \text{step_grad}(\cdot)$
 $\Theta, \text{clusters} \leftarrow \text{detect_cluster}(\cdot)$
 $G \leftarrow \text{gradient}(\cdot)$
end while
return $\Theta, \text{clusters}$

The `gradient` function depends on all the parameters, `step_grad(.)` is just the step part of a gradient descent algorithm : we update the estimation by :

$$\hat{\Theta}_{k+1} \leftarrow \hat{\Theta}_k - h \times \nabla L(\Theta)$$

For the `detect_cluster(.)`, the clusters merged if the distance between two groups C_1 and C_2 is under a small threshold. Then, the coefficient of the new cluster C is computed by the weighted mean of the two other one :

$$\theta_C = \frac{|C_1|\theta_{C_1} + |C_2|\theta_{C_2}}{|C_1| + |C_2|}.$$

We can also try to fuse clusters if the cost function decreases if merging.

4 Clusterpath adapted for HRGM

Now we want to adapt the previous method to the Husler-Reiss graphical models.

4.1 Maximum likelihood for graphical model

For the estimation of the precision matrix for HRGM, (Hentschel, Engelke, and Segers 2023) shows that find the Θ is equivalent to minimise :

$$L(\Theta) = -\log(|\Theta|_+) - \frac{1}{2}tr(\bar{\Gamma}\Theta),$$

where $\bar{\Gamma}$ is an estimation of the variogram matrix Γ and $|\cdot|_+$ the generalised determinant.

For the next, we will first use the l^2 norm penalty and we will try to minimize :

$$L_{\mathcal{P}}(\Theta, \lambda) = L(\Theta) + \lambda \mathcal{P}(\Theta)$$

with $\lambda > 0$ and $\Theta \in \mathcal{P}_d^1$.

4.2 Adaptation of the expression

As $\Theta \in \mathcal{P}_d^1$, there are supplementary conditions on the matrix : the rows must sum to one !

It follows that :

$$a_{kk} = r_{kk} - \sum_{j=1}^K p_j r_{kj}, \quad \forall k \in \{1, \dots, K\}$$

More, we can rewrite the matrix Θ as follow (Touw et al. 2024) :

$$\Theta = \begin{pmatrix} (a_{11} - r_{11})I_{p_1} & 0 & \dots & 0 \\ 0 & (a_{22} - r_{22})I_{p_2} & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & (a_{KK} - r_{KK})I_{p_K} \end{pmatrix} + URU^t,$$

where U is a $d \times K$ matrix such that $u_{ij} = 1$ if $i \in C_j$ and 0 otherwise.

Then we can deduce the computation of the trace of $\bar{\Gamma}\Theta$:

$$\text{tr}(\bar{\Gamma}\Theta) = \text{tr}(\bar{\Gamma}URU^t) + \sum_{k=1}^K (a_{kk} - r_{kk})\text{tr}(\bar{\Gamma}_{C_k}) \left(= \text{tr}(\bar{\Gamma}URU^t) - \sum_{l=1}^K p_l \sum_{k=1}^K r_{kl}\text{tr}(\bar{\Gamma}_{C_k}) \right)$$

with $\bar{\Gamma}_{C_k}$ the $p_k \times p_k$ matrix computed from $\bar{\Gamma}$ with the indices in C_k .

Adaptation of the distance

Let's take square ℓ^2 norm penalty for \mathcal{P} .

For i, j in the same cluster C_k we have :

$$\|\theta_{.i} - \theta_{.j}\|^2 = 2(a_{kk} - r_{kk})^2$$

Here, we wish that for two variables in the same cluster, this distance is equal to zero.

The distance D from (Touw et al. 2024) is built for this reason. But in our case we can upgrade the distance : indeed, our matrix Θ have an additionally constraint that rows sum to 0. Thus, we can remove the $(\theta_{ii} - \theta_{jj})^2$ term as the latter is obviously equal to zero if we are in the same cluster.

Therefore, we will consider the following squared distance for the next :

$$D^2(\theta_{.i}, \theta_{.j}) = \sum_{t \neq i, j} (\theta_{it} - \theta_{jt})^2$$

Now, it is time to write the penalty formula with the R matrix. We have :

- for i, j in the same cluster $D^2(\theta_{.i}, \theta_{.j}) = 0$ (it is the goal of this distance).
- for i, j in respectively the clusters C_k and C_l we have :

$$\begin{aligned} D^2(\theta_{.i}, \theta_{.j}) &= \sum_{t \neq i, j} (\theta_{it} - \theta_{jt})^2 \\ &= \sum_{q \neq k, l} p_m (r_{kq} - r_{lq})^2 + (1 - p_k)(r_{kk} - r_{lk})^2 + (1 - p_l)(r_{ll} - r_{lk})^2 \\ &= \tilde{D}^2(r_{.k}, r_{.l}). \end{aligned}$$

Then, by grouping all the terms in \mathcal{P} , we get :

$$\mathcal{P}(R) = \sum_{l > k} W_{kl} \tilde{D}^2(r_{\cdot k}, r_{\cdot l})$$

with $W_{lk} = \sum_{i \in C_k} \sum_{j \in C_l} w_{ij}$.

Our likelihood can now be expressed as a function of the R -matrix.

4.3 Computation of the derivative

We can see that the penalized negative log-likelihood can be decomposed as a sum of three element :

$$L_{\mathcal{P}}(R, \lambda) = L_{log}(R) + L_{trace}(R) + \lambda L_{pen}(R)$$

So we just need to compute separately all the derivative for each coefficient (in the upper triangular part, by symmetry).

4.3.1 Gradient of $L_{log}(R)$

Let's denote $E = \{(i, j) \in \llbracket 1, d \rrbracket^2, i \leq j\}$ and $F = \{(k, l) \in \llbracket 1, K \rrbracket^2, k \leq l\}$ and we consider \mathbb{R}^E and \mathbb{R}^F the set of real-valued vector indexed by E and F .

One can show that there exists a function $f : \mathbb{R}^F \rightarrow \mathbb{R}^E$ such that :

$$f(R) = U(\Theta)$$

where U is the application defined in (Hentschel, Engelke, and Segers 2023) which maps a matrix to a vector in \mathbb{R}^E containing the entries in the upper triangular part of the matrix.

Then we can defined also the function $g : \mathbb{R}^E \rightarrow \mathbb{R}$ as :

$$g(U(\Theta)) = -\log(|\Theta|_+)$$

which is well defined by symmetry of Θ .

Thus we can see our sub-loglikelihood as :

$$L_{log}(R) = g \circ f(R)$$

And then use the chain rule formula to get the partial derivative for $n \in F$:

$$\frac{\partial L_{log}(R)}{\partial r_n} = \sum_{m \in E} \frac{\partial g}{\partial \theta_m}(f(R)) \times \frac{\partial f_m}{\partial r_n}(R)$$

In the appendix of (Hentschel, Engelke, and Segers 2023), they show that :

$$\frac{\partial g(U(\Theta))}{\partial \theta_m} = \gamma(\Theta^+)_m, \quad \text{for } m \in E$$

Now we just have to compute the derivative of f . Let $m = (i, j) \in E$. We can stress first that :

$$f_m(R) = \begin{cases} r_{kk} - \sum_{l=1}^K p_k r_{kl} & \text{if } i = j \in C_k \\ r_{kl} & \text{if } i \in C_k, j \in C_l, i < j \end{cases}$$

Then we can deduce the expression below :

$$\frac{\partial f}{\partial r_{kl}} = \begin{cases} 1 & \text{if } i \in C_k, j \in C_l, \text{ or } i \in C_l, j \in C_k, \\ -p_k & \text{if } i = j \in C_k \\ -p_l & \text{if } i = j \in C_l \\ 0 & \text{otherwise} \end{cases}$$

We can now associate these expressions to deduce for $k < l$ that :

$$\begin{aligned} \frac{\partial L_{log}(R)}{\partial r_{kl}} &= -\text{tr}[\gamma(\Theta^+)_{C_k}]p_k - \text{tr}[\gamma(\Theta^+)_{C_l}]p_l + \sum_{\substack{i \in C_k \\ i < j}} \sum_{j \in C_l} \gamma(\Theta^+)_{ij} + \sum_{\substack{i \in C_l \\ i < j}} \sum_{j \in C_k} \gamma(\Theta^+)_{ij} \\ &= u_k^t \gamma(\Theta^+) u_l - \text{tr}[\gamma(\Theta^+)_{C_k}]p_k - \text{tr}[\gamma(\Theta^+)_{C_l}]p_l, \end{aligned}$$

with u_k the k -th column of the matrix of clusters U .

Moreover if $k = l$, we get :

$$\begin{aligned} \frac{\partial L_{log}(R)}{\partial r_{kk}} &= \text{tr}[\gamma(\Theta^+)_{C_k}](1 - p_k) + \sum_{\substack{i, j \in C_k \\ i < j}} \gamma(\Theta^+)_{ij} \\ &= \text{tr}[\gamma(\Theta^+)_{C_k}](1 - p_k) + \frac{1}{2} u_k^t \gamma(\Theta^+) u_k \end{aligned}$$

4.3.2 Gradient of $L_{trace}(R)$

We recall that :

$$L_{trace}(R) = \text{tr}(\bar{\Gamma} U R U^t) - \sum_{l=1}^K p_l \sum_{k=1}^K r_{kl} \text{tr}(\bar{\Gamma}_{C_k})$$

As we have for all symmetric matrix A :

$$\frac{\partial \text{tr}(AB)}{\partial b_{ij}} = 2a_{ij}$$

Then we can deduce that for $(k, l) \in F$:

$$\frac{\partial L_{trace}(R)}{\partial r_{kl}} = \begin{cases} 2(U^t \bar{\Gamma} U)_{kk} - p_k \text{tr}(\bar{\Gamma}_{C_k}), & \text{if } k = l \\ 2(U^t \bar{\Gamma} U)_{kl} - p_k \text{tr}(\bar{\Gamma}_{C_l}) - p_l \text{tr}(\bar{\Gamma}_{C_k}) & \text{otherwise.} \end{cases}$$

4.3.3 Gradient of $L_{pen}(R)$

We recall that :

$$L_{pen}(R) = \sum_{l' > k'} W_{k'l'} \tilde{D}^2(r_{\cdot k'}, r_{\cdot l'})$$

So we just need to compute the derivative of \tilde{D}^2 for each coefficient r_{kl} .

Let $(k, l) \in F$.

Thus, for $k < l$, we have :

$$\frac{\partial D^2(r_{\cdot k'}, r_{\cdot l'})}{\partial r_{kl}} = \begin{cases} 2p_k(r_{kl} - r_{kk'}) & \text{if } k \neq k', l = l' \\ 2p_l(r_{kl} - r_{ll'}) & \text{if } k = k', l \neq l' \\ 2(p_k - 1)(r_{kl} - r_{kk}) + 2(p_l - 1)(r_{kl} - r_{ll}) & \text{if } k = k', l = l' \\ 0 & \text{otherwise.} \end{cases}$$

and for $k' < l'$:

$$\frac{\partial D^2(r_{\cdot k'}, r_{\cdot l'})}{\partial r_{kk}} = \begin{cases} 2(p_k - 1)(r_{kk} - r_{kl'}) & \text{if } k = k' \\ 2(p_k - 1)(r_{kk} - r_{k'l'}) & \text{if } k = l' \\ 0 & \text{otherwise.} \end{cases}$$

We can deduce the derivatives :

$$\frac{\partial L_{pen}(R)}{\partial r_{kl}} = \sum_{k'=1}^{l-1} 2W_{k'l} p_k(r_{kl} - r_{kk'}) + \sum_{l'=k+1}^K 2W_{kl'} p_l(r_{kl} - r_{ll'}) + 2W_{kl}(r_{kk} + r_{ll} - 2r_{kl})$$

and :

$$\frac{\partial L_{pen}(R)}{\partial r_{kk}} = 2(p_k - 1) \sum_{q=1}^K W_{qk}(r_{kk} - r_{kq})$$

4.4 Choice of a step for descent

4.5 Simulation study

5 Application on flight delay data

References

- Engelke, Sebastian, and Adrien S. Hitz. 2020. “Graphical Models for Extremes.” *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 82 (4): 871–932. <https://doi.org/10.1111/rssb.12355>.
- Engelke, Sebastian, Alexander Malinowski, Zakhar Kabluchko, and Martin Schlather. 2012. “Estimation of Hüsler-Reiss Distributions and Brown-Resnick Processes.” arXiv. <https://doi.org/10.48550/arXiv.1207.6886>.
- Hentschel, Manuel, Sebastian Engelke, and Johan Segers. 2023. “Statistical Inference for Hüsler-Reiss Graphical Models Through Matrix Completions.” arXiv. <https://doi.org/10.48550/arXiv.2210.14292>.

- Hocking, Toby Dylan, Armand Joulin, Francis Bach, and Jean-Philippe Vert. n.d. “Clusterpath: An Algorithm for Clustering Using Convex Fusion Penalties.”
- Rootzén, Holger, and Nader Tajvidi. 2006. “Multivariate Generalized Pareto Distributions.” *Bernoulli* 12 (5): 917–30. <https://doi.org/10.3150/bj/1161614952>.
- Touw, D. J. W., A. Alfons, P. J. F. Groenen, and I. Wilms. 2024. “Clusterpath Gaussian Graphical Modeling.” arXiv. <https://doi.org/10.48550/arXiv.2407.00644>.